

# Backpropagation Through

# Backpropagation Through **THE VOID**

A person is shown from behind, in a meditative pose with arms raised, reaching towards a bright, glowing light source. The scene is set in a dark, cavernous space with large, curved, organic structures resembling tentacles or cave formations. The overall color palette is dominated by deep blues and blacks, with the light source providing a strong contrast.



# Backpropagation Through **THE VOID**

Will Grathwohl  
Dami Choi  
Yuhuai Wu  
Geoff Roeder  
David Duvenaud




# Where do we see this guy?

$$\mathcal{L}(\theta) = \mathbb{E}_{p(b|\theta)}[f(b)]$$

- Variational Inference
- Hamiltonian Monte Carlo
- Policy Optimization
- Hard Attention

# Bayesian optimization doesn't scale yet



Shahriari et al., 2016

- Bayesopt is usually expensive, relative to model evals
- Global surrogate models not good enough in high dim.
- Even for expensive black-box functions, gradient-based optimization is embarrassingly competitive
- Can we add some cheap model-based optimization to REINFORCE?


# REINFORCE (Williams, 1992)

$$\hat{g}_{\text{REINFORCE}}[f] = f(b) \frac{\partial}{\partial \theta} \log p(b|\theta), \quad b \sim p(b|\theta)$$

- Unbiased
- Works for any  $f$ , not differentiable or even unknown
- high variance



$$\hat{g}_{\text{REINFORCE}}[f] = f(b) \frac{\partial}{\partial \theta} \log p(b|\theta), \quad b \sim p(b|\theta)$$




# Reparameterization Trick:

$$\hat{g}_{\text{reparam}}[f] = \frac{\partial f}{\partial b} \frac{\partial b}{\partial \theta} \quad b = T(\theta, \epsilon), \epsilon \sim p(\epsilon)$$

- Usually lower variance
- Unbiased
- Gold standard, allowed huge continuous models
- Requires  $f(b)$  to be known and differentiable
- Requires  $p(b|\theta)$  to be differentiable



$$\hat{g}_{\text{reparam}}[f] = \frac{\partial f}{\partial b} \frac{\partial b}{\partial \theta} \quad b = T(\theta, \epsilon), \epsilon \sim p(\epsilon)$$



# Concrete/Gumbel-softmax

$$\hat{g}_{\text{concrete}}[f] = \frac{\partial f}{\partial \sigma(z/t)} \frac{\partial \sigma(z/t)}{\partial \theta} \quad z = T(\theta, \epsilon), \epsilon \sim p(\epsilon)$$

- Tune variance vs bias
- Works well in practice for discrete models
- Biased
- $f(b)$  must be known and differentiable
- $p(z|\theta)$  must be differentiable
- Uses undefined behavior of  $f(b)$

# Control Variates

- Allow us to reduce variance of a Monte Carlo estimator

$$\hat{g}_{\text{new}}(b) = \hat{g}(b) - c(b) + \mathbb{E}_{p(b)}[c(b)]$$

- Variance is reduced if  $\text{corr}(g, c) > 0$
- Need to adapt  $g$  as problem changes during optimization


# Our Approach

$$\begin{aligned}\hat{g}_{\text{LAX}} &= g_{\text{REINFORCE}}[f] - g_{\text{REINFORCE}}[c_\phi] + g_{\text{reparam}}[c_\phi] \\ &= [f(b) - c_\phi(b)] \frac{\partial}{\partial \theta} \log p(b|\theta) + \frac{\partial}{\partial \theta} c_\phi(b)\end{aligned}$$



# Our Approach

$$\hat{g}_{\text{LAX}} = \underbrace{g_{\text{REINFORCE}}[f]}_{\text{unbiased}} - \underbrace{g_{\text{REINFORCE}}[c_\phi]}_{\text{biased}} + \underbrace{g_{\text{reparam}}[c_\phi]}_{\text{biased}}$$



# Optimizing the Control Variate

$$\frac{\partial}{\partial \phi} \text{Variance}(\hat{g}) = \mathbb{E} \left[ \frac{\partial}{\partial \phi} \hat{g}^2 \right]$$

- For any unbiased estimator we can get Monte Carlo estimates for the gradient of the variance of  $\hat{g}$
- Use to optimize  $c_\phi$
- Got trick from Ruiz et al. and REBAR paper

# A self-tuning gradient estimator

- Jointly optimize original problem and surrogate together with stochastic optimization
- Requires higher-order derivatives

---

**Algorithm 1** LAX: Optimizing parameters and a gradient control variate simultaneously.

---

**Require:**  $f(\cdot)$ ,  $\log p(b|\theta)$ , reparameterized sampler  $b = T(\theta, \epsilon)$ , neural network  $c_\phi(\cdot)$ ,  
step sizes  $\alpha_1, \alpha_2$

**while** not converged **do**

$\epsilon \sim p(\epsilon)$

▷ Sample noise

$b \leftarrow T(\epsilon, \theta)$

▷ Compute input

$\hat{g}_\theta \leftarrow [f(b) - c_\phi(b)] \nabla_\theta \log p(b|\theta) + \nabla_\theta c_\phi(b)$

▷ Estimate gradient of objective

$\hat{g}_\phi \leftarrow \partial \hat{g}_\theta^2 / \partial \phi$

▷ Estimate gradient of variance of gradient

$\theta \leftarrow \theta - \alpha_1 \hat{g}_\theta$

▷ Update parameters

$\phi \leftarrow \phi - \alpha_2 \hat{g}_\phi$


▷ Update control variate

**end while**

**return**  $\theta$

---

$$\hat{g}_{\text{LAX}} = [f(b) - c_{\phi}(b)] \frac{\partial}{\partial \theta} \log p(b|\theta) + \frac{\partial}{\partial \theta} c_{\phi}(b)$$





**What about discrete  
variables?**

# Extension to discrete $p(b|\theta)$

$$\hat{g}_{\text{DLAX}} = f(b) \frac{\partial}{\partial \theta} \log p(b|\theta) - c_\phi(z) \frac{\partial}{\partial \theta} \log p(z|\theta) + \frac{\partial}{\partial \theta} c_\phi(z)$$

$$b = H(z), z \sim p(z|\theta)$$

$$H(z) = b \sim p(b|\theta)$$

- Unbiased for all  $c_\phi$

# Extension to discrete $p(b|\theta)$

$$\hat{g}_{\text{RELAX}} = [f(b) - c_\phi(\tilde{z})] \frac{\partial}{\partial \theta} \log p(b|\theta) + \frac{\partial}{\partial \theta} c_\phi(z) - \frac{\partial}{\partial \theta} c_\phi(\tilde{z})$$

$$b = H(z), z \sim p(z|\theta), \tilde{z} \sim p(z|b, \theta)$$

- Main trick introduced in REBAR (Tucker et al. 2017).
- We just noticed it works for any  $c()$
- REBAR is special case of RELAX where  $c$  is concrete relaxation
- We use autodiff to tune entire surrogate, not just temperature



# Toy Example

$$\mathbb{E}_{p(b|\theta)}[(t - b)^2]$$

- Used to validate REBAR (used  $t = .45$ )
- We use  $t = .499$
- REBAR, REINFORCE extremely slow in this case
- Can RELAX improve?




# Toy Example



- massively reduced variance
- Surrogate needs time to catch up

# Analyzing the Surrogate

- REBAR's fixed surrogate only adapts temperature param.
- RELAX surrogate balances REINFORCE variance and reparameterization variance
- Optimal surrogate is always smooth



# Define functions, not computation graphs

```
def relax(params, est_params, noise_u, noise_v, func_vals):
    samples = bernoulli_sample(params, noise_u)
    log_temperature, nn_params = est_params

    def surrogate(relaxed_samples):
        return nn_predict(nn_params, relaxed_samples)

    def surrogate_cond(params):
        cond_noise = conditional_noise(params, samples, noise_v) # z tilde
        return concrete(params, log_temperature, cond_noise, surrogate)

    grad_surrogate = elementwise_grad(concrete)(params, log_temperature, noise_u, surrogate)
    surrogate_cond, grad_surrogate_cond = value_and_grad(surrogate_cond)(params)
    return reinforce(params, noise_u, func_vals - surrogate_cond) + \
        grad_surrogate - grad_surrogate_cond

def relax_all(params, est_params, noise_u, noise_v, f):
    # Returns objective, gradients, and gradients of variance of gradients.
    func_vals = f(bernoulli_sample(params, noise_u))
    var_vjp, grads = make_vjp(relax, argnum=1)(params, est_params, noise_u, noise_v, func_vals)
    d_var_d_est = var_vjp(2 * grads / grads.shape[0])
    return func_vals, grads, d_var_d_est
```

# Discrete VAEs

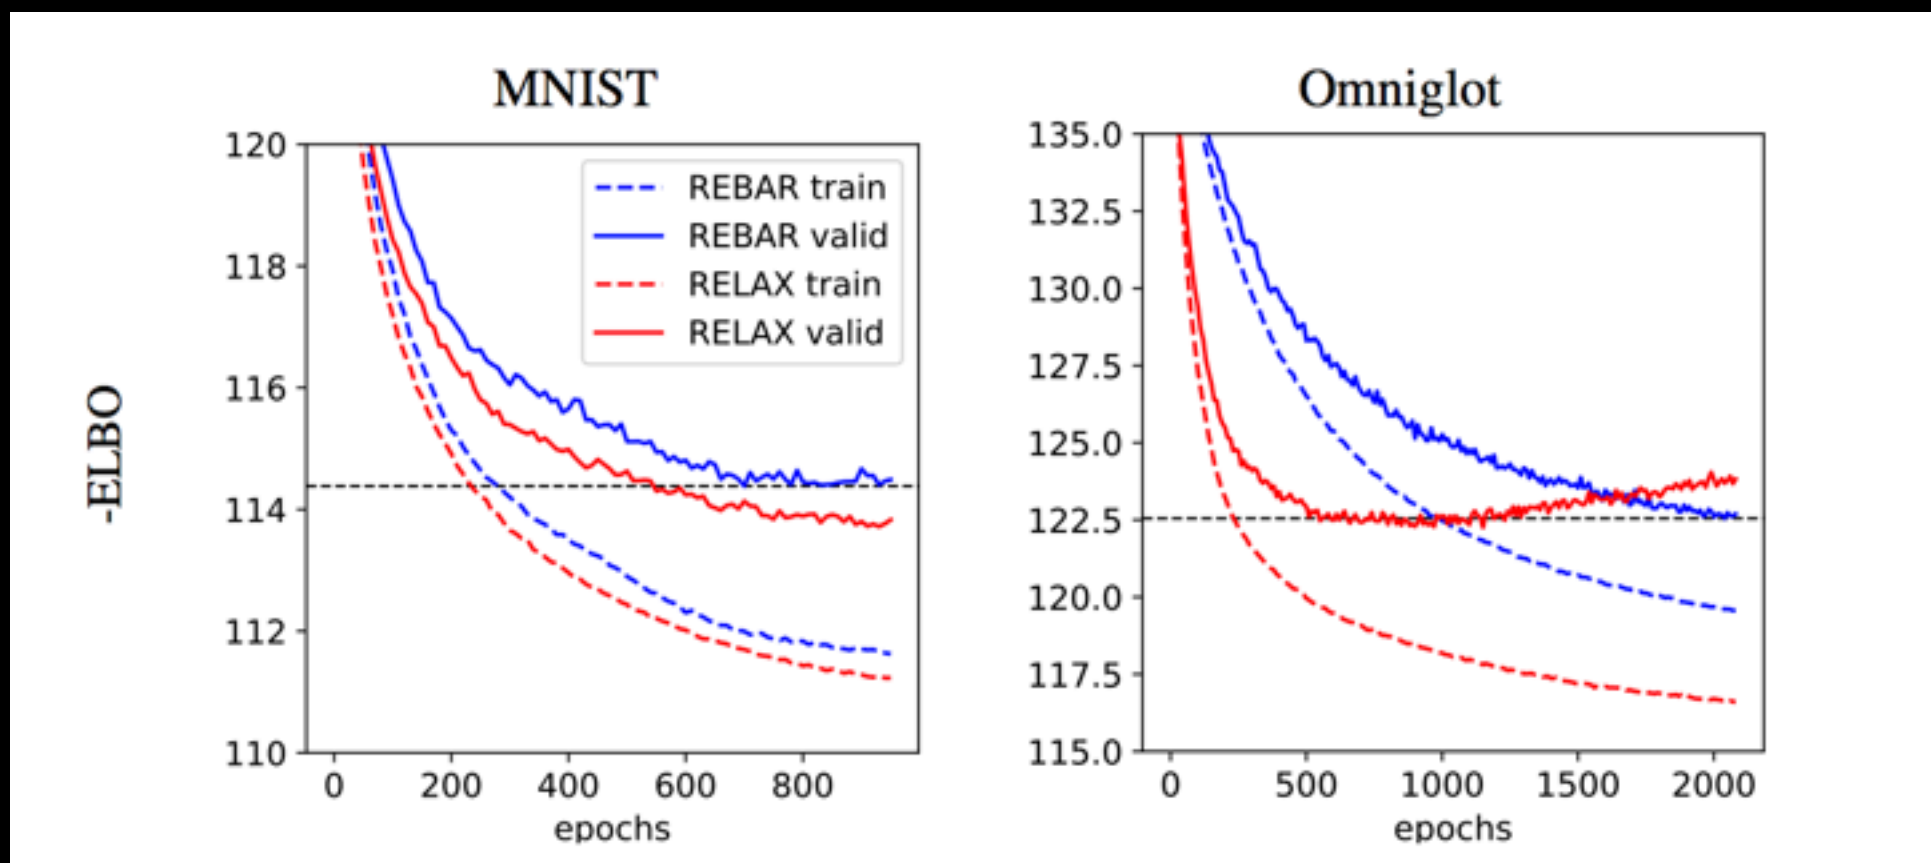
$$\log p(x) \geq \mathcal{L}(\theta) = \mathbb{E}_{q(b|x)} [\log p(x|b) + \log p(b) - \log q(b|x)]$$

- Latent state is 200 Bernoulli variables
- Can't use reparameterization trick
- Can still use our knowledge of structure of model, combining REBAR and RELAX:

$$c_{\phi}(z) = f(\sigma_{\lambda}(z)) + r_{\rho}(z)$$



# Bernoulli VAE Results



Dataset	Model	Concrete	NVIL	MuProp	REBAR	RELAX
<b>MNIST</b>	Nonlinear	-102.2	-101.5	-101.1	-81.01	<b>-78.13</b>
	linear one-layer	-111.3	-112.5	-111.7	-111.6	<b>-111.20</b>
	linear two-layer	-99.62	-99.6	-99.07	-98.22	<b>-98.00</b>
<b>Omniglot</b>	Nonlinear	-110.4	-109.58	-108.72	-56.76	<b>-56.12</b>
	linear one-layer	-117.23	-117.44	-117.09	-116.63	<b>-116.57</b>
	linear two-layer	-109.95	-109.98	-109.55	-108.71	<b>-108.54</b>

Table 1: Highest training ELBO for discrete variational autoencoders.

# Rederiving Actor-Critic

$$\hat{g}_{AC} = \sum_{t=1}^T \frac{\partial \log \pi(a_t | s_t, \theta)}{\partial \theta} \left[ \sum_{t'=t}^T r_{t'} - c_{\phi}(s_t) \right]$$

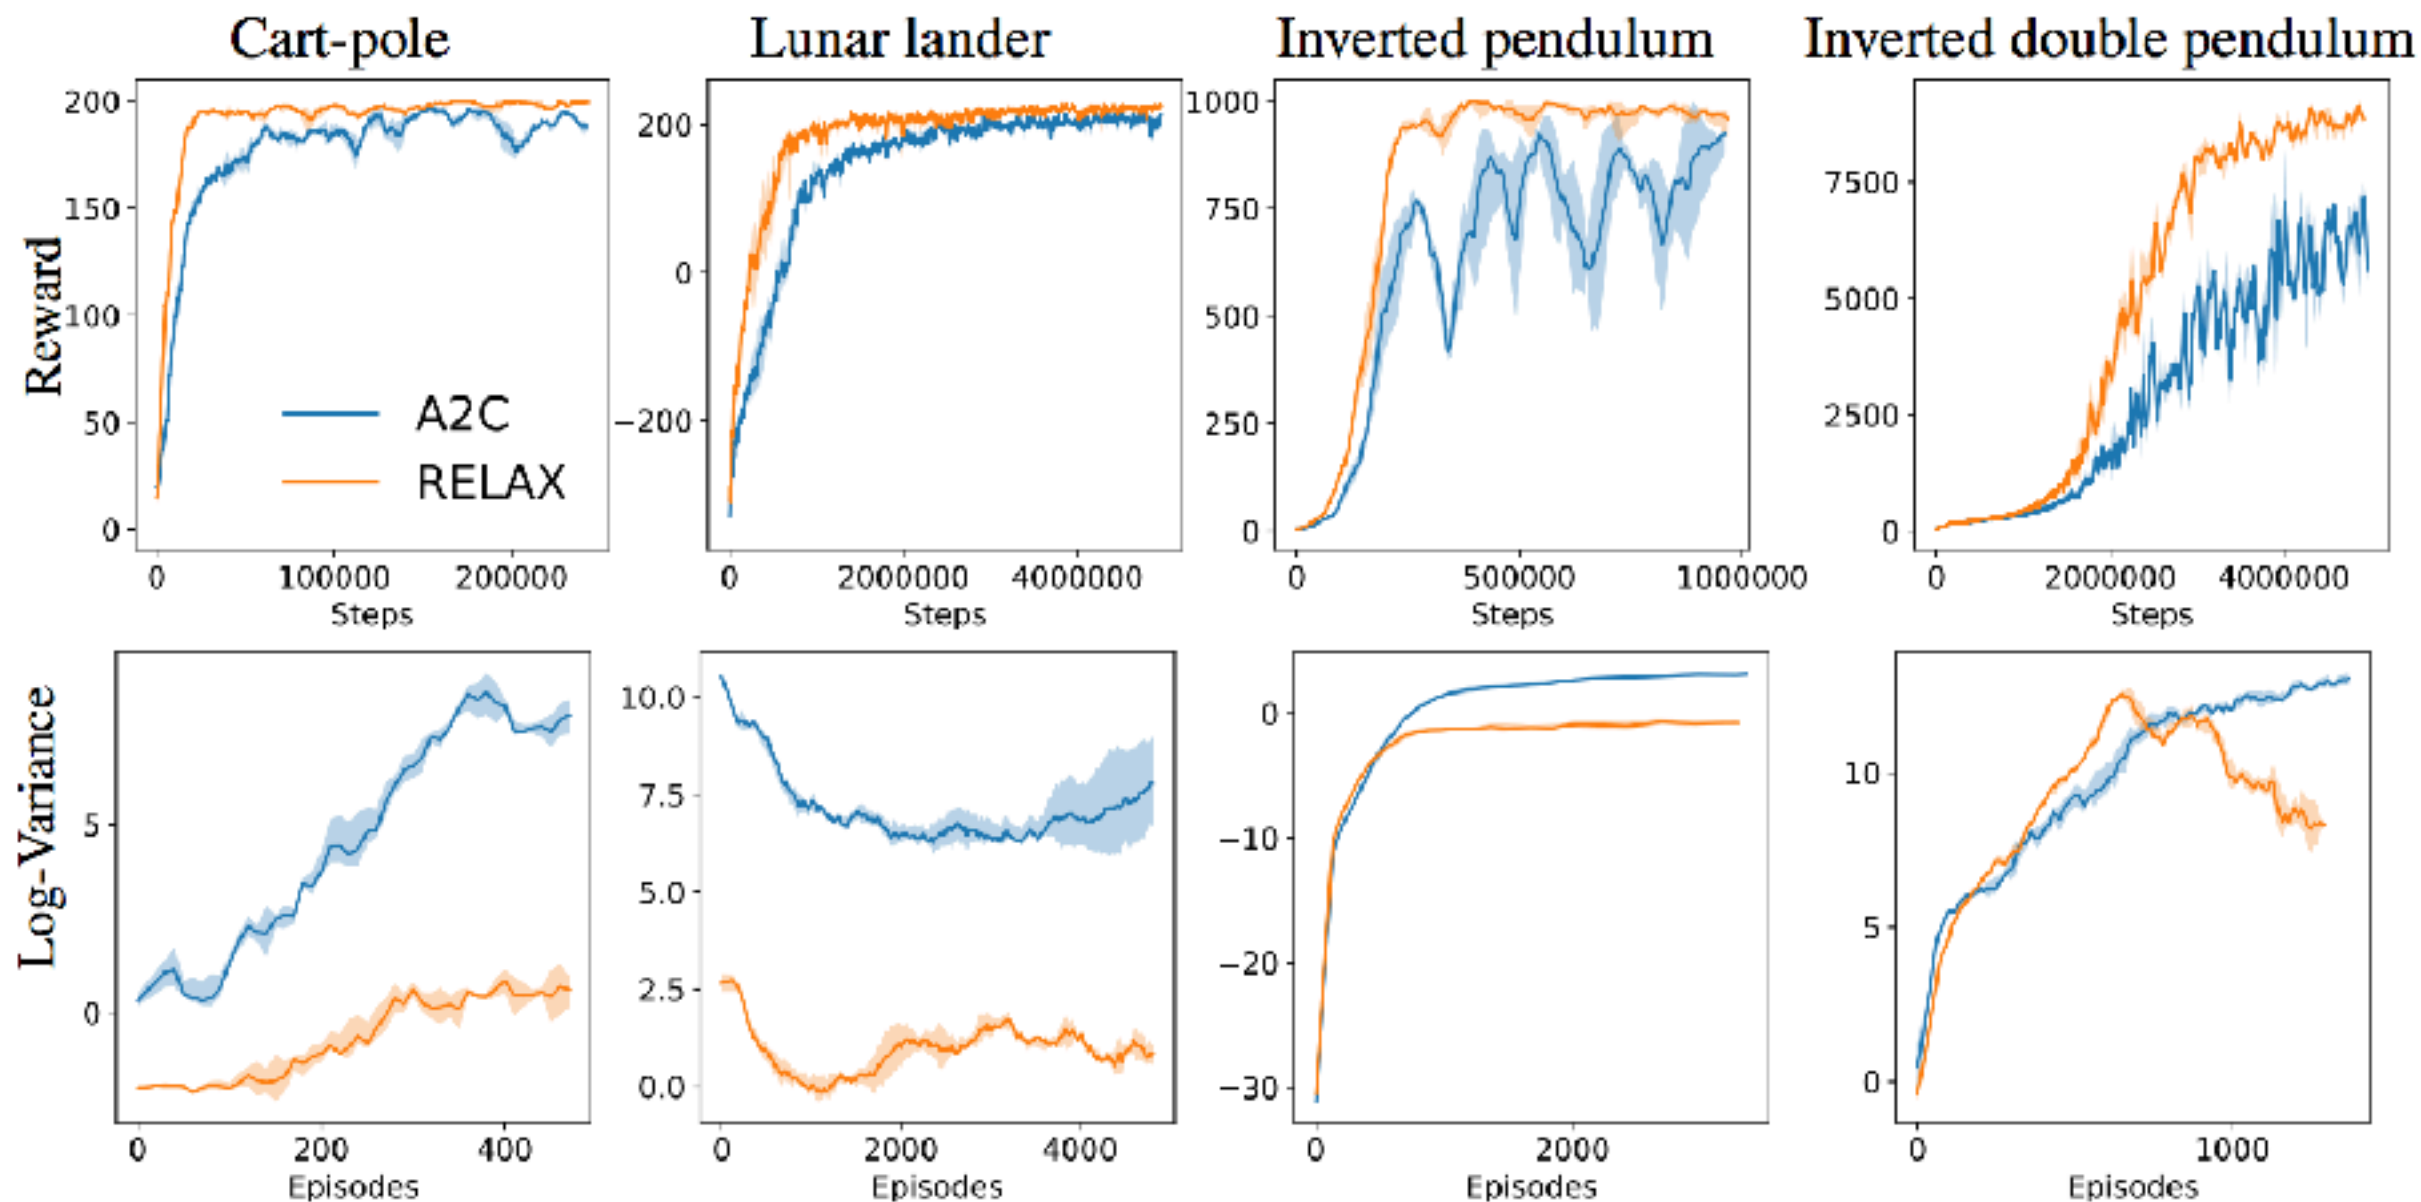
- $c_{\phi}$  is an estimate of the value function
- This is exactly the REINFORCE estimator using an estimate of the value function as a control variate
- Why not use action in control variate?
- Dependence on action would add bias

# LAX for RL

$$\hat{g}_{\text{LAX}} = \sum_{t=1}^T \frac{\partial \log \pi(a_t | s_t, \theta)}{\partial \theta} \left[ \sum_{t'=t}^T r_{t'} - c_{\phi}(s_t, a_t) \right] + \frac{\partial}{\partial \theta} c_{\phi}(s_t, a_t)$$

- Action-dependence in control variate
- unbiased for policy, and unbiased for baseline
- Standard baseline optimization methods minimize squared error from reward or value function. We directly minimize variance.

# Model-Free RL “Results”



- Faster convergence, but real story is unbiased critic updates.
- Excellent criticism of experimental setup in “The Mirage of Action-Dependent Baselines in Reinforcement Learning” (Tucker et al. 2018). Better experiments would examine high-dimensional action regime.

# RELAX Properties


- Pros:

- unbiased
- low variance (after tuning)
- usable when  $f(b)$  is unknown, or not differentiable
- useable when  $p(b|\theta)$  is discrete

- Cons:

- need to define surrogate
- when progress is made, need to wait for surrogate to adapt
- Higher-order derivatives still awkward in TF and PyTorch

# Local surrogates are a nice compromise



Shahriari et al., 2016

- Global surrogate models not good enough in high dim.
- Local surrogates are less demanding to construct
- Local minima less of a problem in high dim.
- Want low variance? Take gradient of variance



# Aside: Evolution Strategies optimize a linear surrogate

$$\hat{w} = (X^T X)^{-1} X^T y$$


$$\approx \mathbb{E} [(X^T X)^{-1} X^T y]$$

$$= [I\sigma^2]^{-1} X^T y$$

$$= [I\sigma^2]^{-1} (\epsilon\sigma)^T y$$

$$= \sum_i \frac{\epsilon_i y_i}{\sigma}$$

$$= \sum_i \frac{\epsilon_i f(\epsilon_i \sigma)}{\sigma}$$


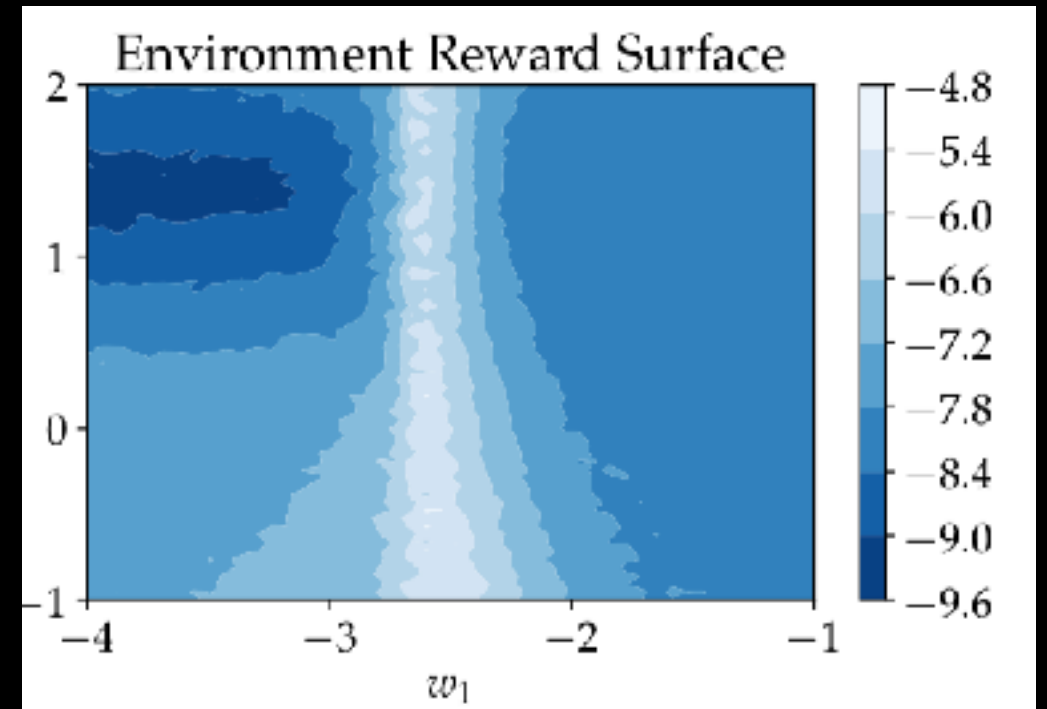


$$\epsilon \sim \mathcal{N}(0, I)$$

$$x = \epsilon\sigma$$

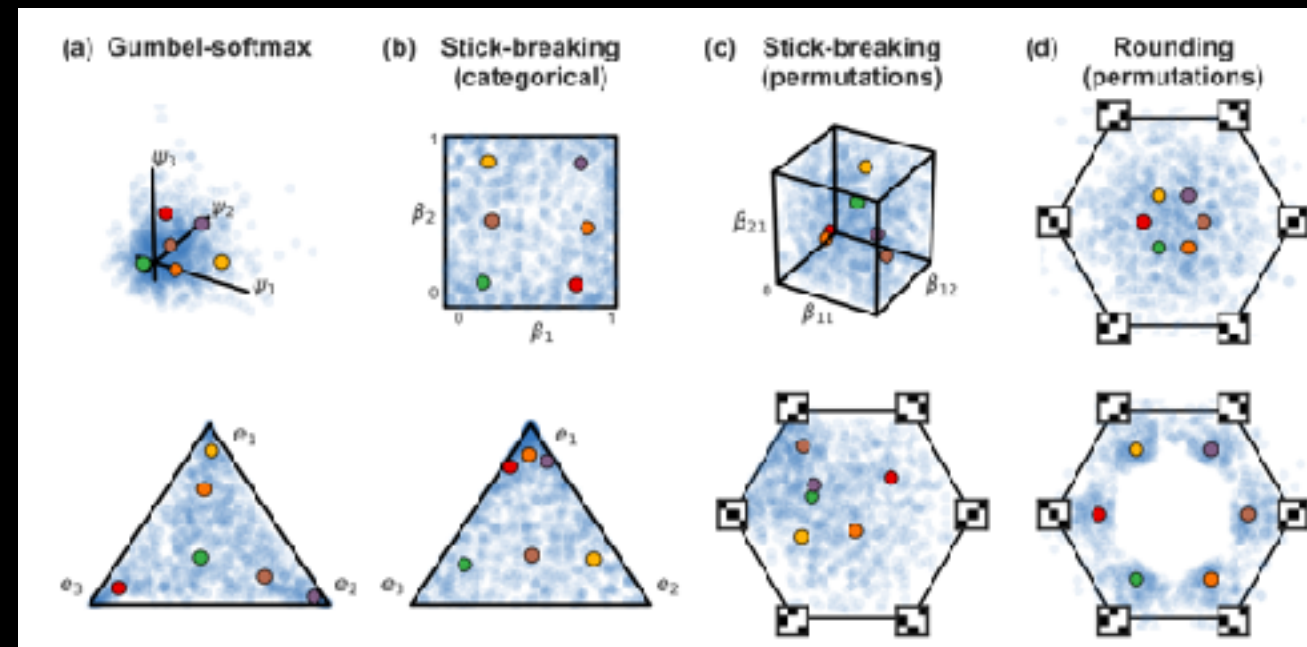
# Aside: Evolution Strategies optimize a linear surrogate

- Throws away all observations each step
- Use a neural net surrogate, and experience replay
- Distributed ES algorithm works for any gradient-free optimization algorithm
- w/ students Geoff Roeder, Yuhuai (Tony) Wu, Jaiming Song




# Future Work

- What does the optimal surrogate look like? Use calculus of variations.
- Train the surrogate off-policy:  $c_\phi(a, s, \pi)$
- REBAR, RELAX for more complicated discrete objects, e.g. trees
- Mostly open problem: Control variates for sequential discrete choices.




Linderman et al., 2017



Mena et al., 2018





Will Grathwohl, Dami Choi, Yuhuai Wu,  
Geoffrey Roeder, Jesse Bettencourt,  
David Duvenaud



# Thanks!

<https://github.com/duvenaud/relax>



# Speed Comparison

